

Gradients and backpropagation: where to start when a loss curve fails Worksheet

Trace one prediction to a parameter update, then use tiny-batch tests, input checks, and training logs to diagnose NaNs, flat loss, and false improvement.

01	Reader decision and user	
02	Input, data boundary and prohibited action	
03	Simple baseline and acceptable output	
04	Fixed cases, failure and handoff	
05	What someone can check in the system: version, test, trace, or result	
06	What still needs explaining before work use: owner, cost, risk, and next step	